

Assembly: load/store operations; arithmetic and bitwise operations; translating assembly code to low-level C code

COSC 208, Introduction to Computer Systems, 2022-03-22

Announcements

- Project 2 due Thursday, Mar 31

Outline

- Warm-up
- Assembly
- Load/store operations
- Arithmetic and bitwise operations
- Translating assembly code to low-level C code

Before break

Language forms

0000	0000	0000	0000	0000	0000	0000	0000
0000	0000	0000	0000	0000	0000	0000	0000
9222	9120	1121	A120	1121	A121	7211	0000
0000	0001	0002	0003	0004	0005	0006	0007
0008	0009	000A	000B	000C	000D	000E	000F
0000	0000	0000	FE10	FACE	CAFE	ACED	CEDE
1234	5678	9ABC	DEF0	0000	0000	F00D	0000
0000	0000	EEEE	1111	EEEE	1111	0000	0000
B1B2	F1F5	0000	0000	0000	0000	0000	0000

Machine language

- Not portable
 - Specific to hardware
- Simple
 - Each instruction does a simple task – poor ratio of functionality to code size
- Not human readable
 - Not structured
 - Requires lots of effort!
 - Requires tool support

We need assembly languages!

```

        mov     w1, 0
loop:   cmp     w0, 1
        ble     endloop
        add     w0, w0, #1
        ands    wzr, w0, #1
        beq     else
        add     w2, w0, w0
        add     w0, w0, w2
        add     w0, w0, 1
        b       endif
else:   asr     w0, w0, 1
endif:  b       loop
endloop:

```

```

count = 0;
while (n>1)
{
    count++;
    if (n&1)
        n = n*3+1;
    else
        n = n/2;
}

```

High-level language: C code

- Portable
- To varying degrees
- Complex
- One statement can do much work – good ratio of functionality to code size
- Human readable
- Structured – `if()`, `for()`, `while()`, etc.

Assembly Languages

- Not portable
 - Each assembly lang instruction maps to one machine lang instruction
- Simple
 - Each instruction does a simple task
- Human readable (In the same sense that Polish is human readable, if you know Polish.)

Why learn enough assembly?

Knowing assembly language helps you:

- Write faster code
 - In assembly language
 - In a high-level language!
- Write safer code
 - Understanding mechanism of potential security problems helps you avoid them – even in high-level languages
- Understand what's happening *under the hood*
 - Someone needs to develop future computer systems
 - Intuition; some part for OS are written in assembly
- Become more comfortable with levels of abstraction
 - Become a better programmer!

Assembly Architecture Types

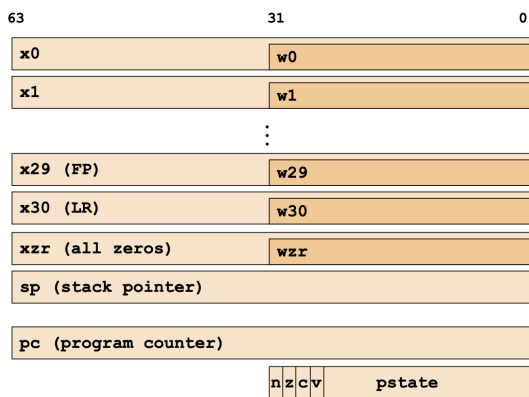
- Instruction set architectures (ISAs) --- defines the low-level instructions a central processing unit (CPU), or processor, can execute
 - Most popular Complex Instruction Set Computer (CISC) architecture: x86 (includes IA32 and x86-64)
 - Most popular Reduced Instruction Set Computer (RISC) architecture: ARM
 - Other common RISC architectures: SPARC (Scalable Processor Architecture), MIPS (Microprocessor without Interlocked Pipelined Stages), PowerPC, ARC (Argonaut RISC Core)

New

Operands

- Registers
 - General purpose: `w0` through `w30` (32-bit) and `x0` through `x30` (64-bit)
 - Stack pointer (top of current stack frame): `sp`
- Constant -- e.g., `#0x20` vs `#12` (hexa vs decimal)
- Memory
 - Dereference --- e.g., `[x1]`
 - Add to (offset from) memory address, then dereference --- e.g., `[sp, #16]`

Registers



Load/store operations

- Typical pattern:
 - Load data from RAM to registers
 - Manipulate data in registers
 - Store data from registers to RAM

On RISC, this pattern is enforced

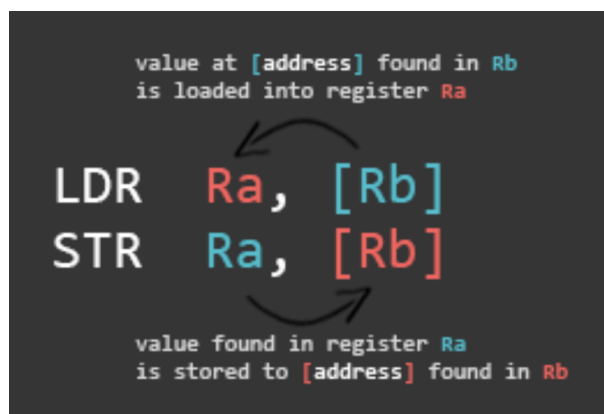
- Manipulation instructions can *only* access registers
- Known as a **Load/store architecture**

CISC does more at once, compact instructions, but slower. Tradeoff

LDR is used to load something from memory into a register, and

STR is used to store something from a register to a memory address. Specifically,

- LDR operation: loads the value at the [memory address] found in `Rb` to the destination register `Ra`.
- STR operation: stores the value found in `Ra` to the [memory address] found in `Rb`.



- Load (**ldr**) — copies a value from a specific location in main memory into a specific register on the CPU
 - Can copy either 64-bits or 32-bits at a time, depending on whether the instruction specifies **x** or **w** register, respectively
 - Similar to dereferencing a pointer and storing the value in a local variable
 - For example, if we treat registers as if they were variable names, **ldr x0, [x1]** is similar to **x0 = *x1**, where **x0** has type **long** and **x1** has type **long ***
- Store (**str**) — copies a value from a specific register on the CPU into a specific location in main memory
 - Again, can copy 64-bits or 32-bits at a time, depending on whether the instruction specifies **x** or **w** register, respectively
 - Similar to storing the value in a local variable at a memory location pointed-to by a pointer
 - For example, if we treat registers as if they were variable names, **str x0, [x1]** is similar to ***x1 = x0**, where **x0** has type **long** and **x1** has type **long ***

Warm-up: Load/store operations

- Q1: Write the C code equivalent for each line of assembly, treating registers as if they were variable names.

- **ldr x0, [sp]**

```
x0 = *sp
```

- **str w0, [sp]**

```
*sp = w0
```

- **ldr x1, [sp,#12]**

```
x1 = *(sp + 12)
```

- **str x2, [x3,#0x10]**

```
*(x3 + 0x10) = x2
```

Arithmetic and logical operations

- (Most) arithmetic and bitwise operations take three operands
 - Register into which to store the result of the operation
 - Register containing the first operand
 - Constant value or register containing the second operand
- Q2: Write the C code equivalent for each line of assembly, treating registers as if they were variable names.

- `lsl w0, w1, w2`

```
w0 = w1 << w2  
  
//equivalent to 2^n
```

- `and w3, w4, #20`

```
w3 = w4 & 20
```

- `mul w5, w6, #0x11`

```
w5 = w6 * 0x11
```

- `sdiv x7, x8, x9`

```
x7 = x8 / x9  
//signed integer division used 64-bit value, hence x registers
```

The `udiv` and `sdiv` instructions operate on 32-bit and 64-bit data respectively. Note that you cannot multiply 32-bit registers with 64 bit registers.

Mapping assembly code to C code

- Q3: The following C code was compiled into assembly.

```
1  #include <stdio.h>
2  int years_to_double(int rate) {
3      int ruleof72 = 72;
4      int years = ruleof72 / rate;
5      return years;
6  }
7  int main() {
8      int r = 10;
9      int y = years_to_double(r);
10     printf("With an interest rate of %d%% it will take ~%d
           years to double your money\n", r, y);
11 }
```

For each line of assembly, indicate which original line of C code (above) the assembly instruction was derived from.

```
000000000000076c <years_to_double>:
76c: d10083ff  sub sp, sp, #0x20    // 2
770: b9000fe0  str w0, [sp, #12]    // 2
774: 52800900  mov w0, #0x48        // 3
778: b9001be0  str w0, [sp, #24]    // 3
77c: b9401be1  ldr w1, [sp, #24]    // 4
780: b9400fe0  ldr w0, [sp, #12]    // 4
784: 1ac00c20  sdiv w0, w1, w0      // 4
788: b9001fe0  str w0, [sp, #28]    // 4
78c: b9401fe0  ldr w0, [sp, #28]    // 5
790: 910083ff  add sp, sp, #0x20    // 5
794: d65f03c0  ret                  // 5
```

- What do each of the columns contain?
 - Code memory address
 - Bytes corresponding to instruction
 - Operation
 - Operands
- Viewing assembly code
 - Compile: `gcc -fomit-frame-pointer -o years years.c`
 - Disassemble executable: `objdump -d years > years.txt`
- Assembly code (excerpt from `years.txt`)

Translating assembly code to low-level C code

The following C program (`operands.c`) has been compiled into assembly:

```
int operandsA(int a) {
    return a;
}

long operandsB(long b) {
    return b;
}

int operandsC(int *c) {
    return *c;
}

long operandsD(long *d) {
    return *d;
}

int main() {
    operandsA(5);
    operandsB(5);
    int x = 5;
    operandsC(&x);
    long y = 5;
    operandsD(&y);
}
```

- Q4: Write the C code equivalent for each line of assembly, treating registers as if they were variable names. The assembly code for the `operandsA` function has already been translated into C code.

```
00000000000007ec <operandsA>:
7ec: d10043ff sub sp, sp, #0x10 // sp = sp - 0x10
7f0: b9000fe0 str w0, [sp, #12] // *(sp + 12) = w0
7f4: b9400fe0 ldr w0, [sp, #12] // w0 = *(sp + 12)
7f8: 910043ff add sp, sp, #0x10 // sp = sp + 0x10
7fc: d65f03c0 ret // return

0000000000000800 <operandsB>:
800: d10043ff sub sp, sp, #0x10 // sp = sp - 0x10
804: f90007e0 str x0, [sp, #8] // *(sp + 8) = x0
808: f94007e0 ldr x0, [sp, #8] // x0 = *(sp + 8)
80c: 910043ff add sp, sp, #0x10 // sp = sp + 0x10
810: d65f03c0 ret // return

0000000000000814 <operandsC>:
814: d10043ff sub sp, sp, #0x10 // sp = sp - 0x10
818: f90007e0 str x0, [sp, #8] // *(sp + 8) = x0
81c: f94007e0 ldr x0, [sp, #8] // x0 = *(sp + 8)
820: b9400000 ldr w0, [x0] // w0 = *x0
824: 910043ff add sp, sp, #0x10 // sp = sp + 0x10
828: d65f03c0 ret // return

000000000000082c <operandsD>:
82c: d10043ff sub sp, sp, #0x10 // sp = sp - 0x10
830: f90007e0 str x0, [sp, #8] // *(sp + 8) = x0
834: f94007e0 ldr x0, [sp, #8] // x0 = *(sp + 8)
838: f9400000 ldr x0, [x0] // x0 = *x0
```

83c:	910043ff	add sp, sp, #0x10	// sp = sp + 0x10
840:	d65f03c0	ret	// return

- Q5: How does the assembly code for *operandsA* and *operandsB* differ? Why?
 - *operandsA* takes and returns an int, which is 32-bits, whereas *operandsB* takes and returns a long, which is 64-bits, so:
 - *operandsA* uses *w0* while *operandsB* uses *x0*
 - *operandsA* stores the parameter at *sp + 12* while *operandsB* stores the parameter at *sp + 8*
- Q6: How does the assembly code for *operandsB* and *operandsD* differ? Why?
 - *operandsB* takes and returns a long, whereas *operandsD* takes a pointer to a long and returns a long, so:
 - *operandsD* must dereference the pointer (*ldr x0, [x0]*) before returning
- Q7: How does the assembly code for *operandsC* and *operandsD* differ? Why?
 - *operandsC* takes a pointer to an int and returns an int, whereas *operandsD* takes a pointer to a long and returns a long
 - both take a memory address (a 64-bit value), which is initially in *x0* and stored at *sp + 8*
 - the dereference of the pointer is a 32-bit value in *operandsC* and a 64-bit value in *operandsD*, so the value is loaded into *w0* in *operandsC* and *x0* in *operandsD*