Multiprocessing: scheduling

COSC 208, Introduction to Computer Systems, 2022-04-26

Outline

- Scheduling processes
- First In First Out (FIFO) scheduling
- Shortest Job First (SJF) scheduling
- Preemption
- Shortest Time-to-Completion First (STCF) scheduling
- Round Robin (RR) scheduling

Scheduling processes

- OS decides which process to run and for how long
- *What factors should the OS consider when making these decisions?*
  - Time for process to complete
  - Overhead of context switching
  - Fairness
  - User interaction
- For now, consider one of these metrics: turnaround time
  - $T_{turnaround} = T_{complete} - T_{arrive}$
- For now, assume a process starts and runs to completion—i.e., no I/O and no preemption

First In First Out (FIFO) scheduling

Process	Arrival time	Duration
A	0	15
B	5	15
C	10	15

- *What is the average turnaround time for the processes using FIFO?*

- *What happens if A's duration is 60?*

- *What happens if C's duration is 60?*

- *How can we change the schedule so the average turnaround time when A's duration is 60 is more like the average turnaround time when C's duration is 60?*

Shortest Job First (SJF) scheduling

Process	Arrival time	Duration
A	0	60
B	0	15
C	0	15

- *What is the average turnaround time for the above processes using SJF?*

Process	Arrival time	Duration
A	0	60
B	5	15
C	10	15

- *What is the average turnaround time for the above processes using SJF?*

- We're back to FIFO—*What happened!? How can we fix this?*

Preemption

- OS only regains control when a system call occurs—e.g., read/write file, yield
    - Syscalls may occur infrequently, or never, due to program design, bugs, or malicious behavior
- *How does an OS forcibly regain control?*—set a timer that raises an interrupt
    - Interrupt causes a trap instruction to be executed
    - Interrupts can also be raised by devices—e.g., Network Interface Card (NIC)
- *What must the OS do if it decides to run another process?*—perform a context switch
    - Save the machine state associated with the process that was running—in particular, the contents of all registers are saved in the process's control structure
    - Restore the machine state associated with the process that should run—again, the contents of all registers are loaded from the process's control structure

Shortest Time-to-Completion First (STCF) scheduling

- Allow preemption
- If a process arrives that has less computation remaining than the currently running process, then preempt the current process and run the new process
- Also known as Preemptive Shortest Job First (PSJF)

Process	Arrival time	Duration
A	0	60
B	5	15
C	10	15

- *What is the average turnaround time for the above processes using STCF scheduling?*

- Now consider response time: $T_{response} = T_{first\_run} - T_{arrive}$
- *What is the average response time for the same processes using STCF scheduling?*
- Now consider wait time: $T_{wait} = \sum (T_{start\_run} - T_{become\_ready})$

- *What major assumption have we made thus far that is impractical in a real system?*

Round Robin (RR)

- Let each process run for a small amount of time, then switch to the next process; when you get to the last process, then start again with the first process and repeat

Process	Arrival time	Duration
A	Just before 0	60
B	Just before 5	15
C	Just before 10	15

- *What is the average turnaround time and response time for the above processes assuming we let a process run for 5 seconds before switching processes?*

  - Average Turnaround =
  - Average Response =
  - Average Wait =
- In practice, there is a queue of processes that are in the ready state, resulting in the following schedule:

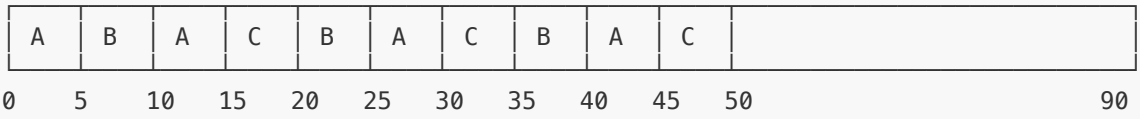

- *Determine the schedule for the above process with a time quantum of 10.*

  - Average Turnaround =
  - Average Response =
  - Average Wait =
- *What happens to average response time as we increase the time quantum?*